

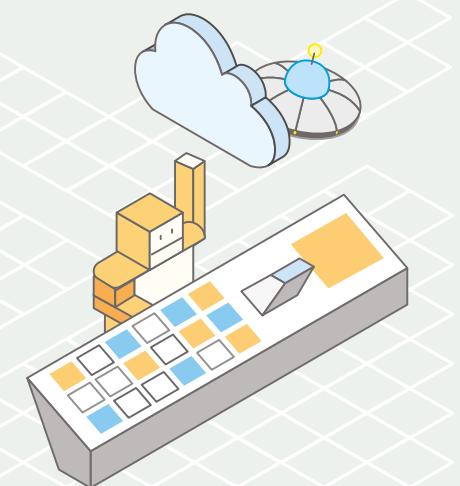
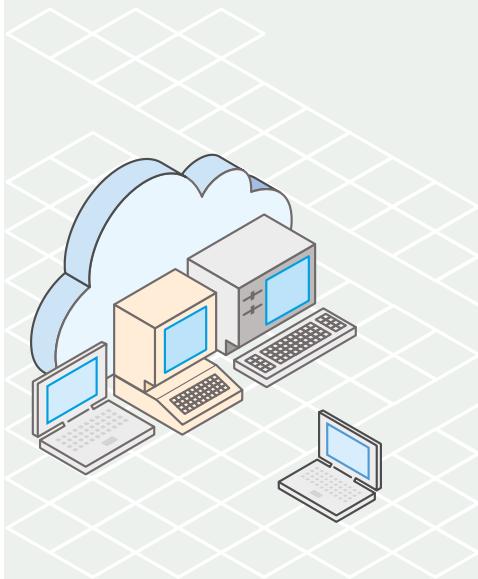


Lambda@Edge

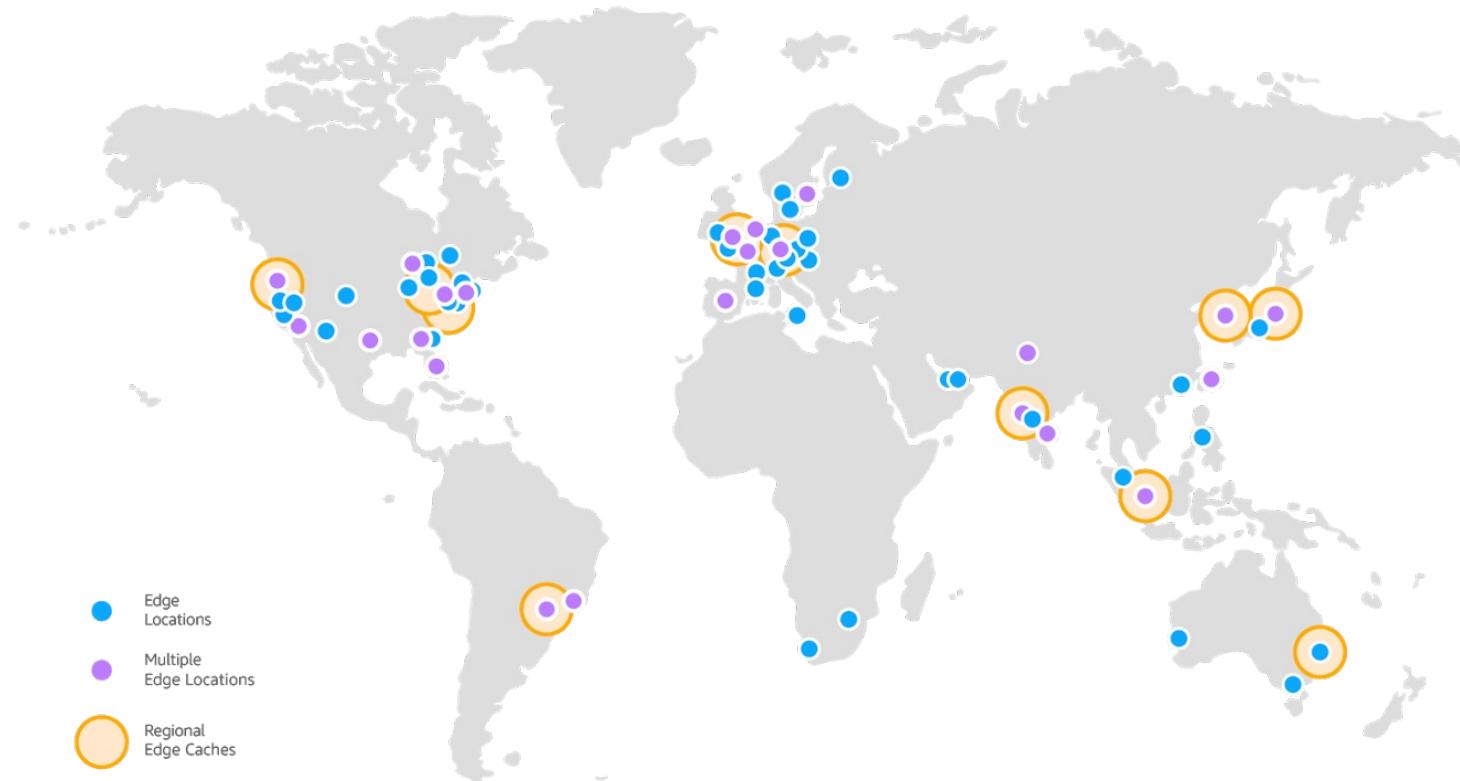
Customize your content delivery

Achraf Souk, Solutions Architect – Edge Services

17/10/2018



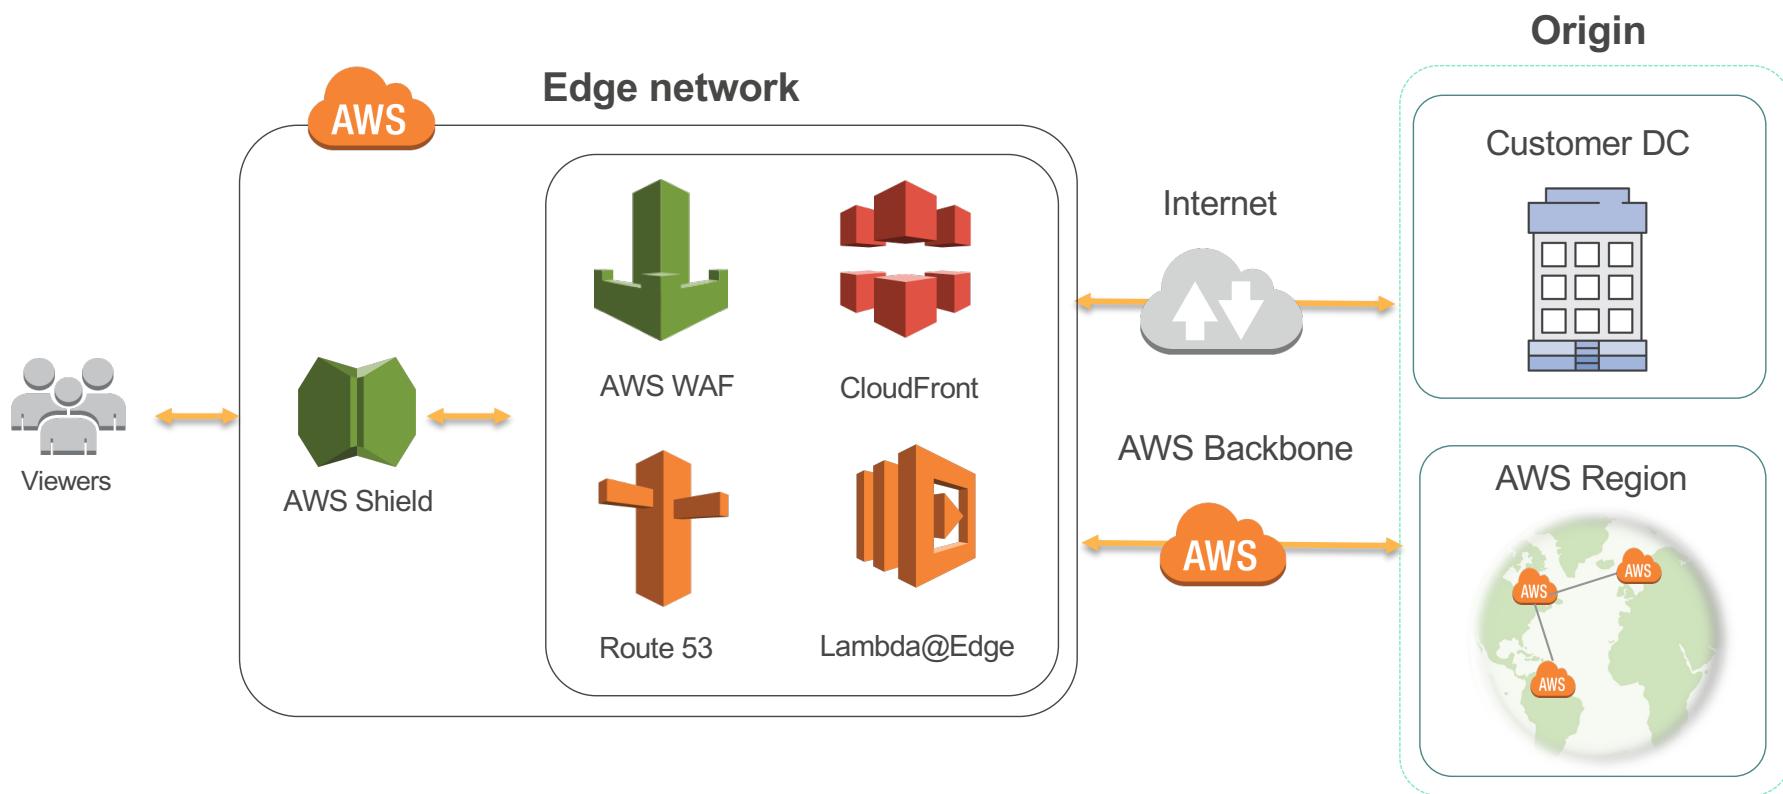
AWS Edge Network



**138 PoPs
63 cities
29 countries**



Edge Services



In the Nordics

- 3 PoPs in Stockholm, Sweden
- 1 PoP in Helsinki, Finland
- 1 PoP in Oslo, Norway
- 1 PoP in Copenhagen, Denmark

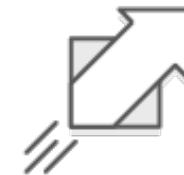
... and a region in Stockholm
coming soon!



Lambda is serverless



No servers to provision
or manage



Scales with usage



Never pay for idle

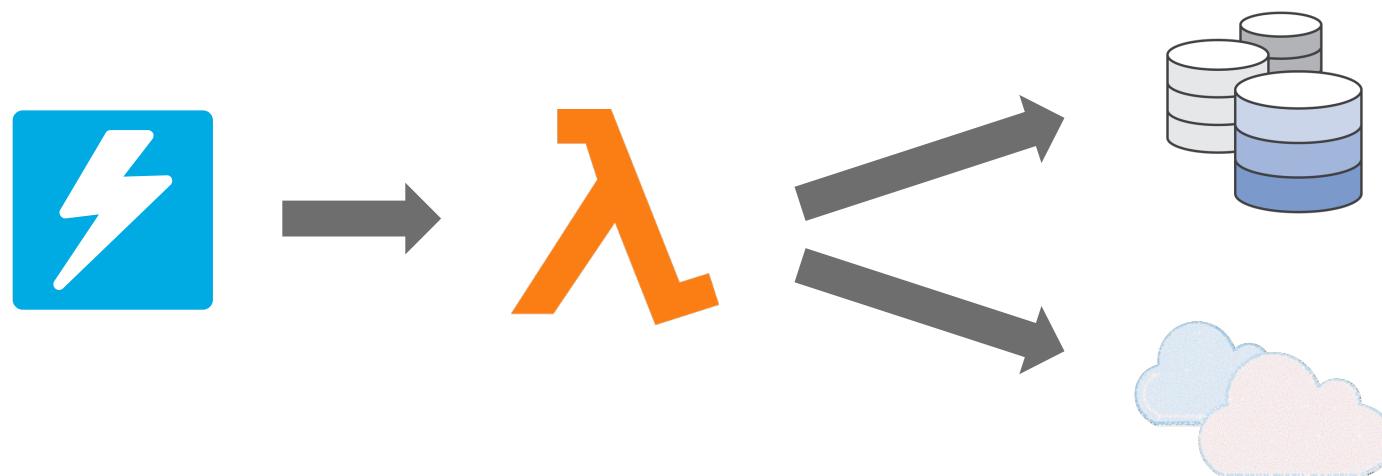


Built-in availability
and fault tolerance



How it works

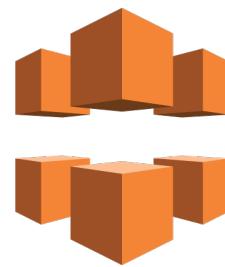
EVENT SOURCE FUNCTION SERVICES



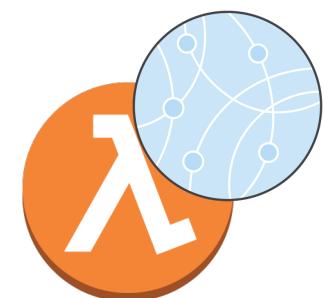
Lambda@Edge is distributed Lambda



AWS Lambda



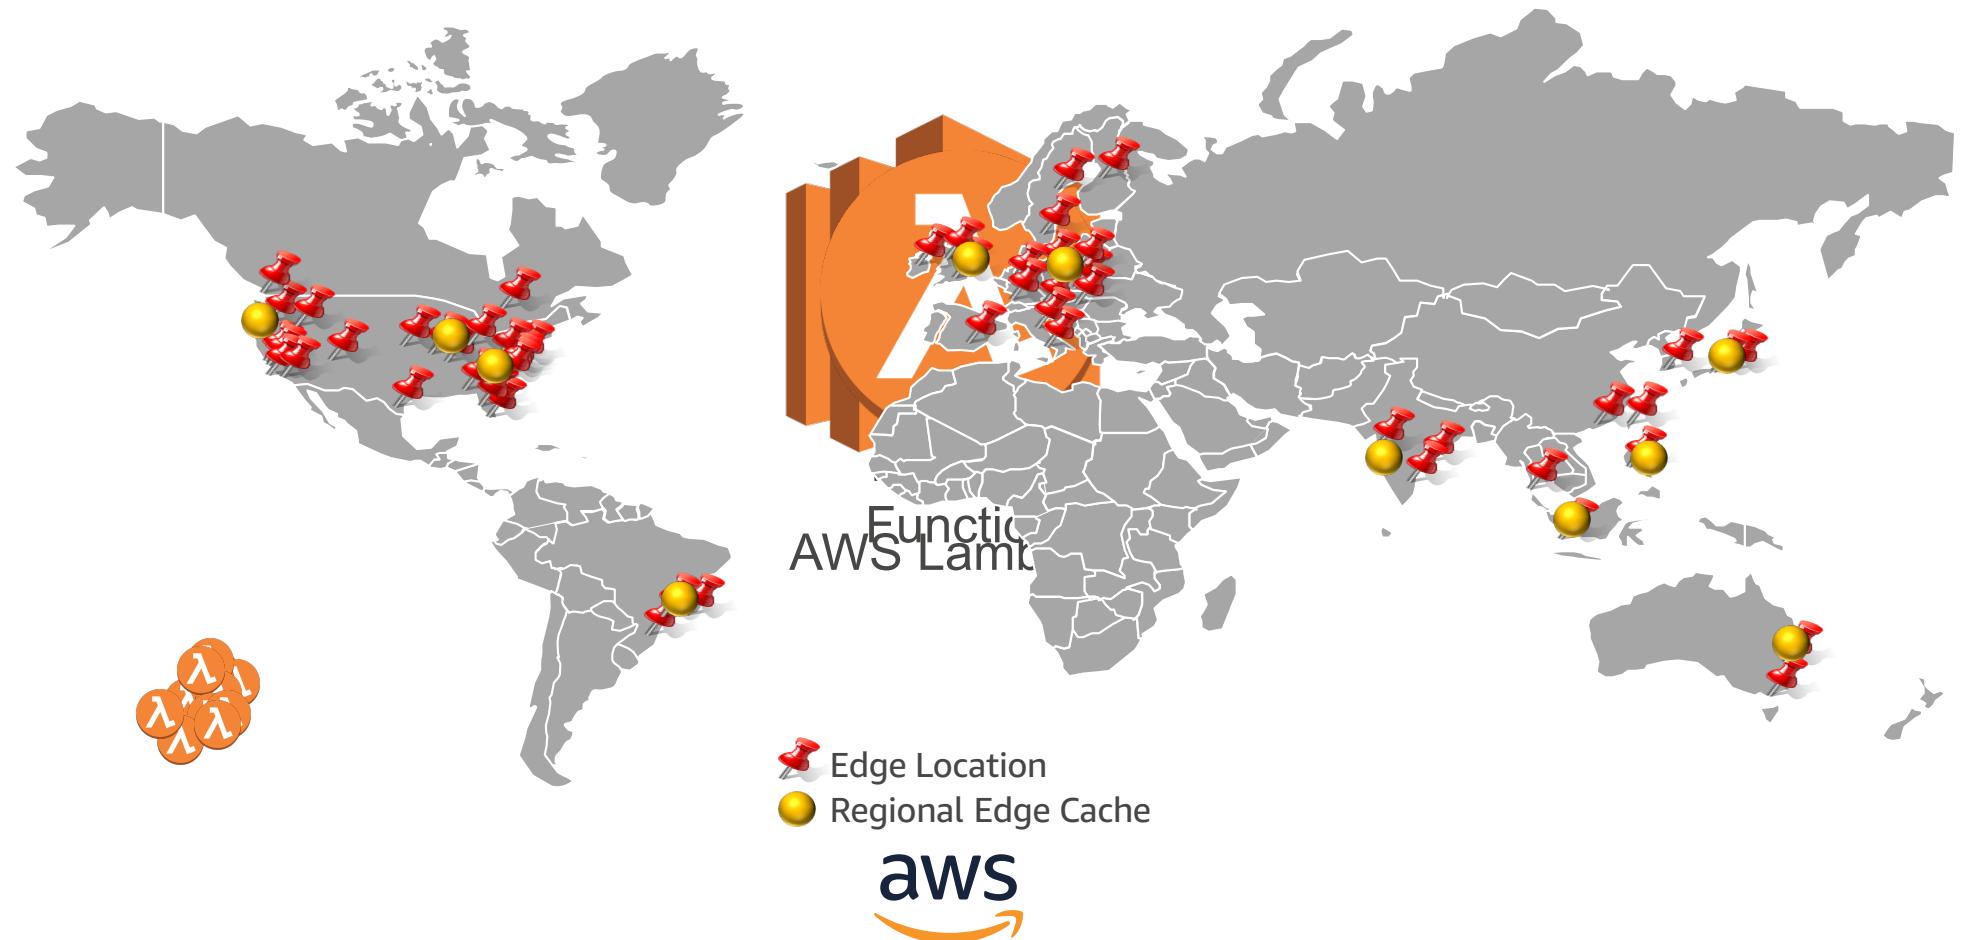
Amazon CloudFront



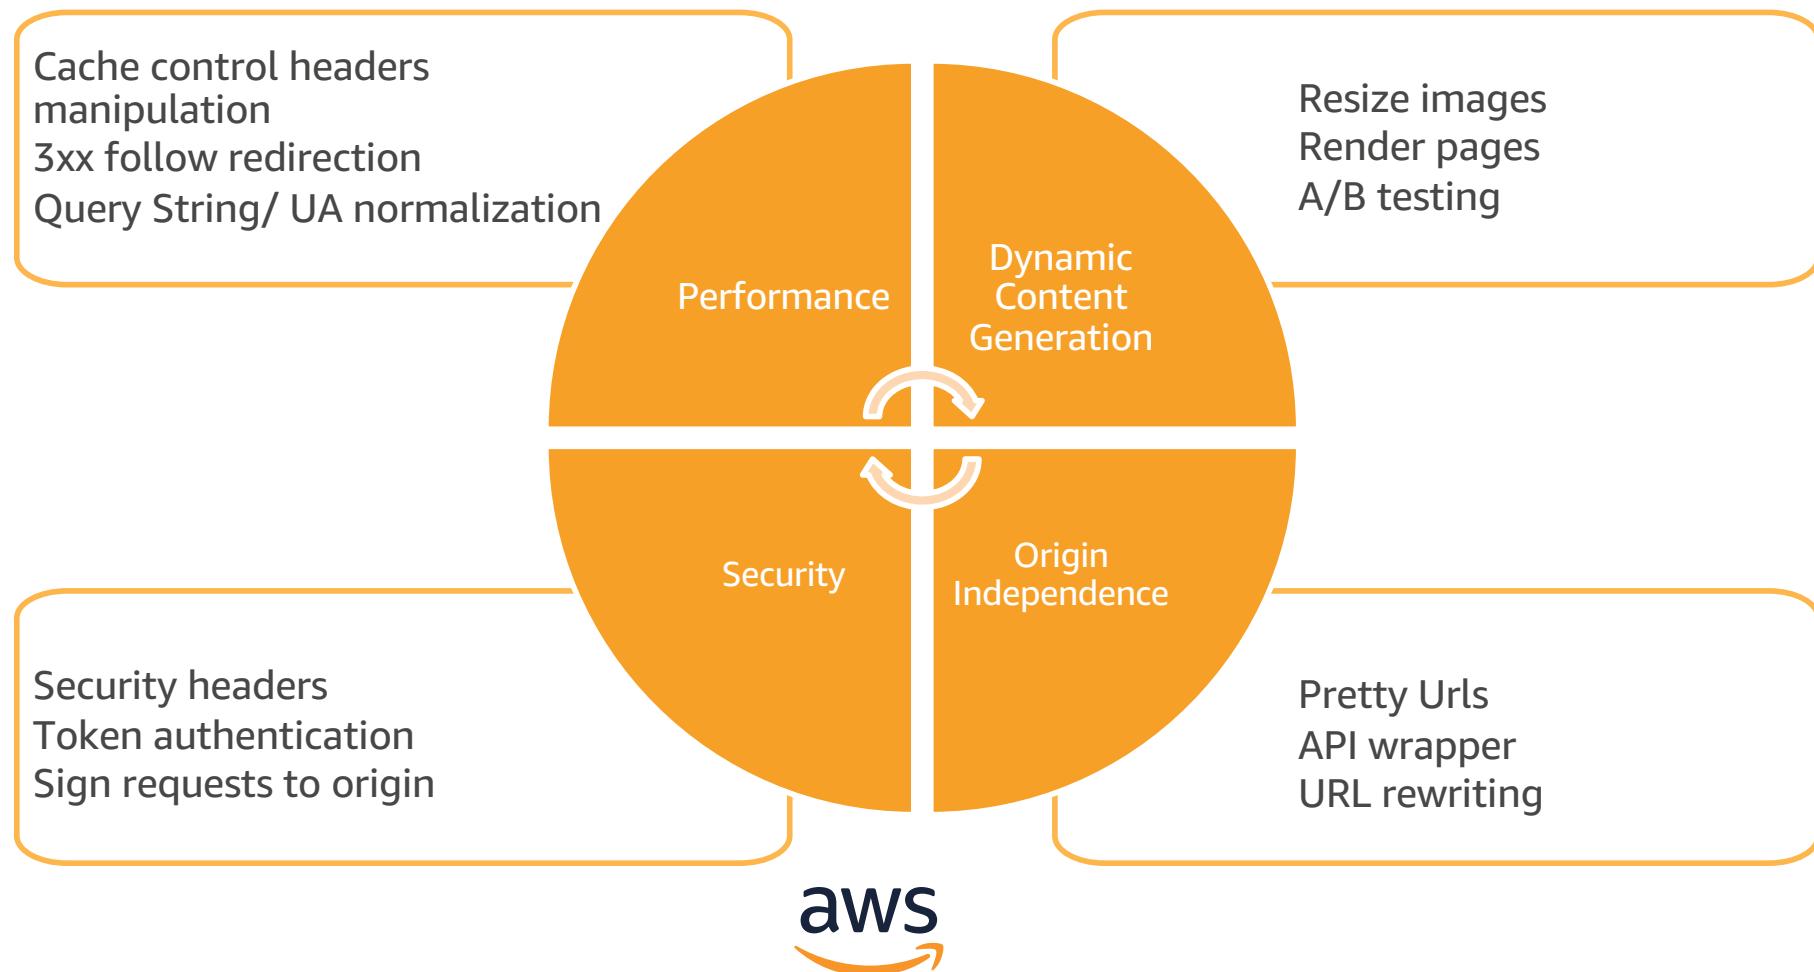
Lambda@Edge



Write once, run Lambda functions globally



Use cases



Bot management - Datadome

- DataDome is an SaaS cybersecurity solution for web and mobile applications that analyzes and manages non-human traffic in real time.
- DataDome uses Lambda@Edge to **simplify** the onboarding process for their customers, and to make their service **available globally** AWS edge locations.



Video streaming - Blockbuster

- Blockbuster provides transactional video-on-demand services to viewers in the nordics with hundreds of thousands of subscribers in Denmark.
- Blockbuster used Lambda@Edge to **solve Smart TV compatibility issues**, by rewriting the manifest to a version understandable by these devices.



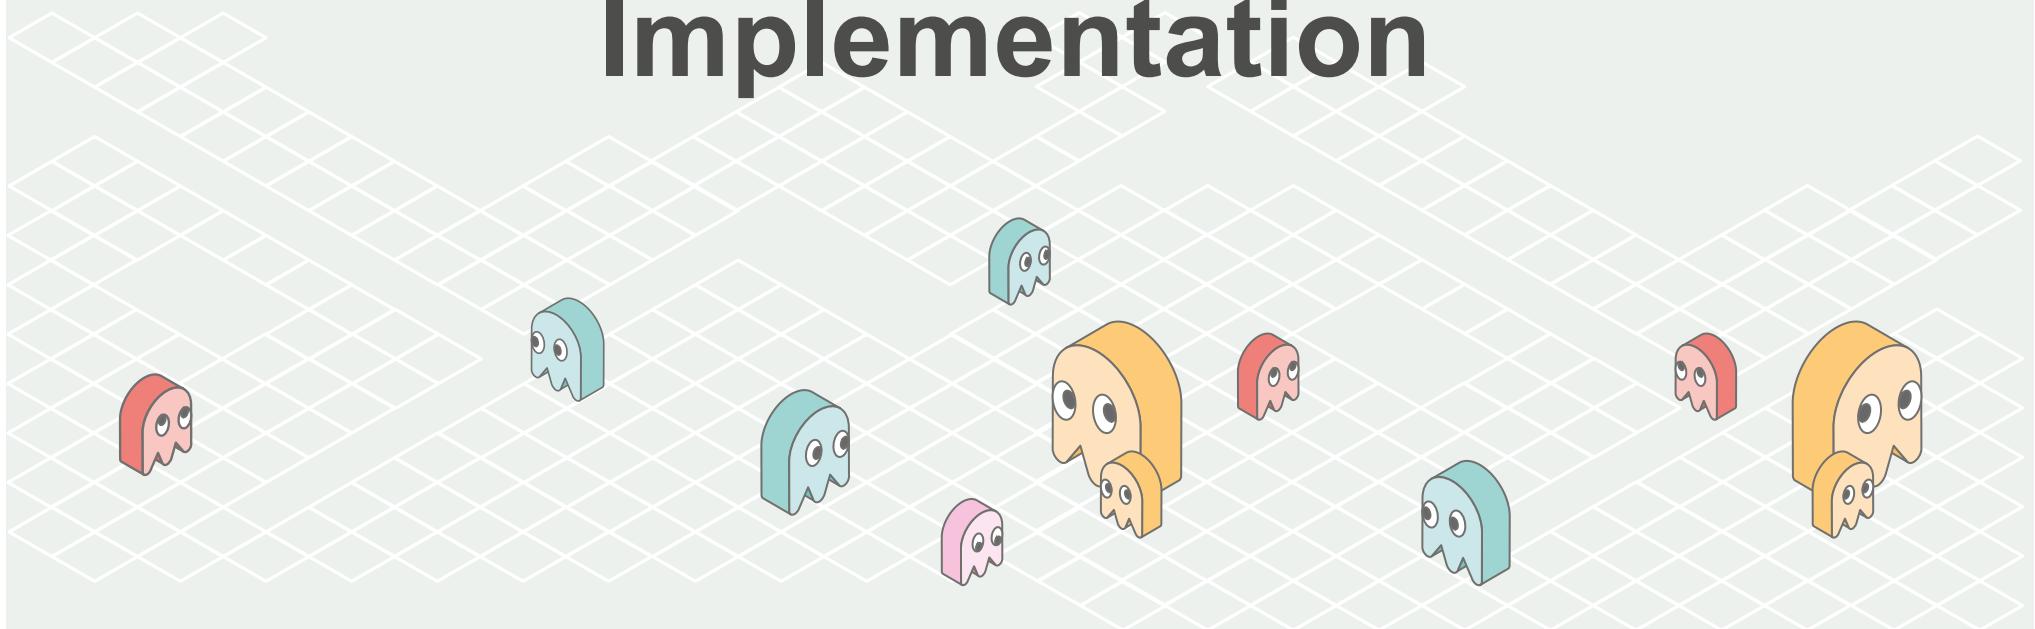
On AWS Blog

- Leveraging Lambda@Edge for **AdTech**: Cookie Syncing at the Edge
- Visitor Prioritization on **e-Commerce** Websites with CloudFront and Lambda@Edge
- Serving **Private Content** Using Amazon CloudFront & AWS Lambda@Edge
- Building a Serverless Subscription Service using Lambda@Edge for **Publishers**

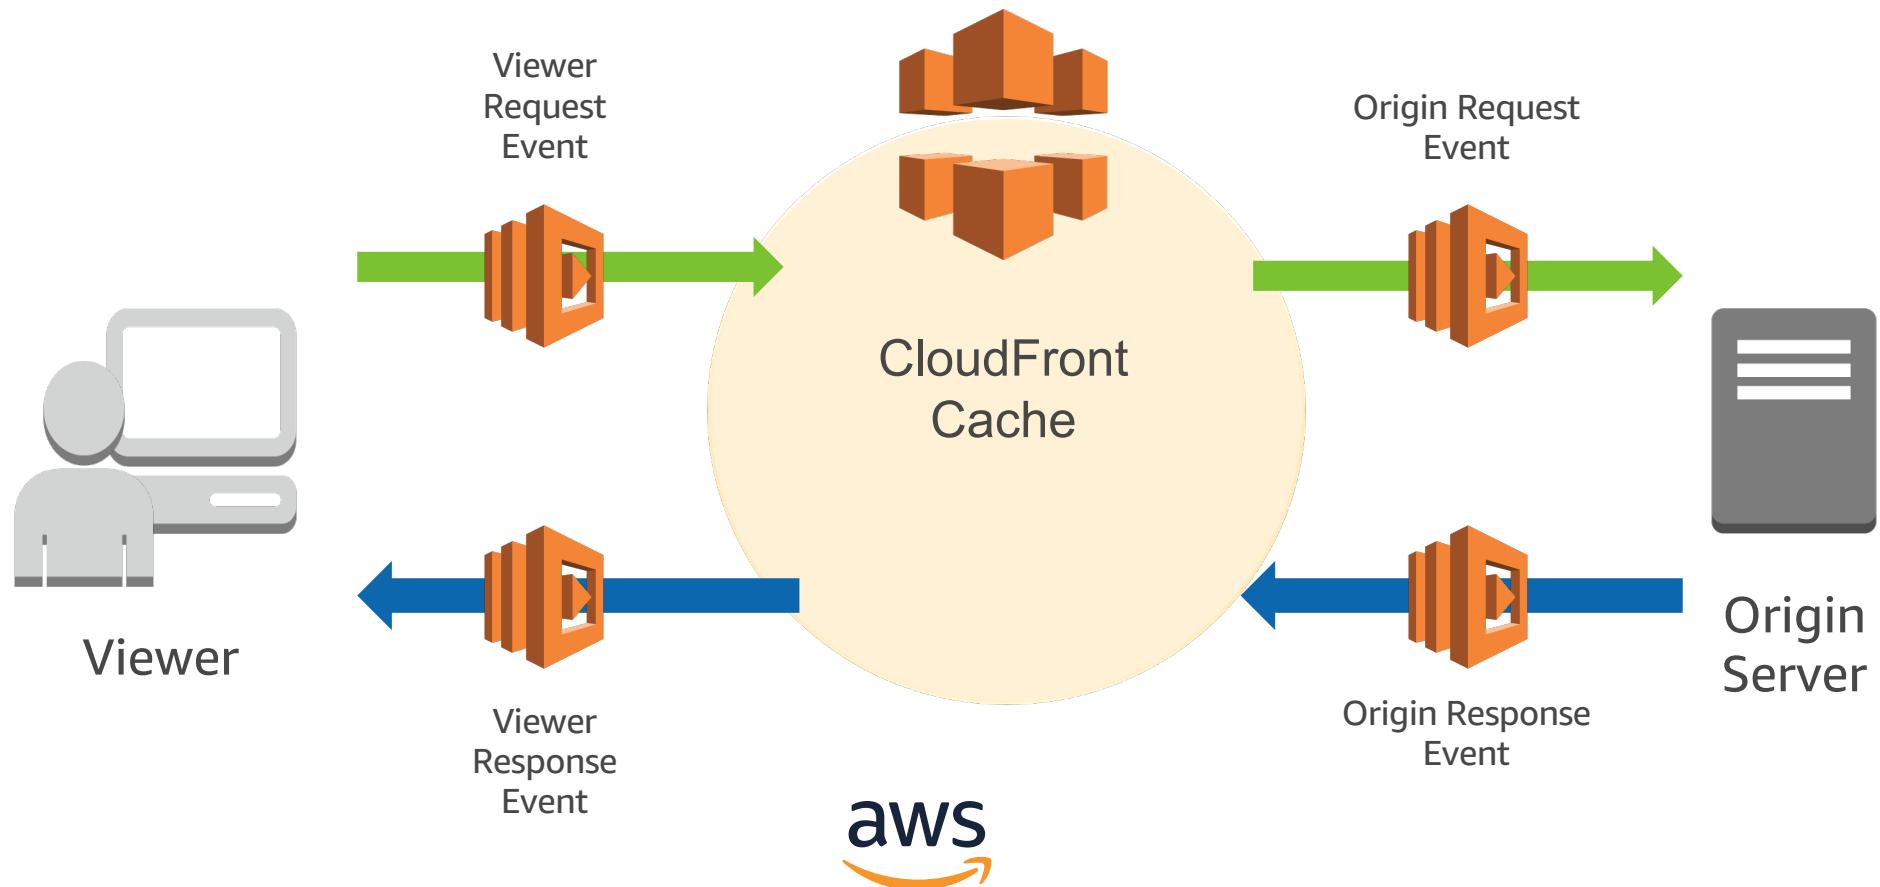




Implementation



CloudFront triggers



Using Lambda@Edge

Author in
N.Virginia

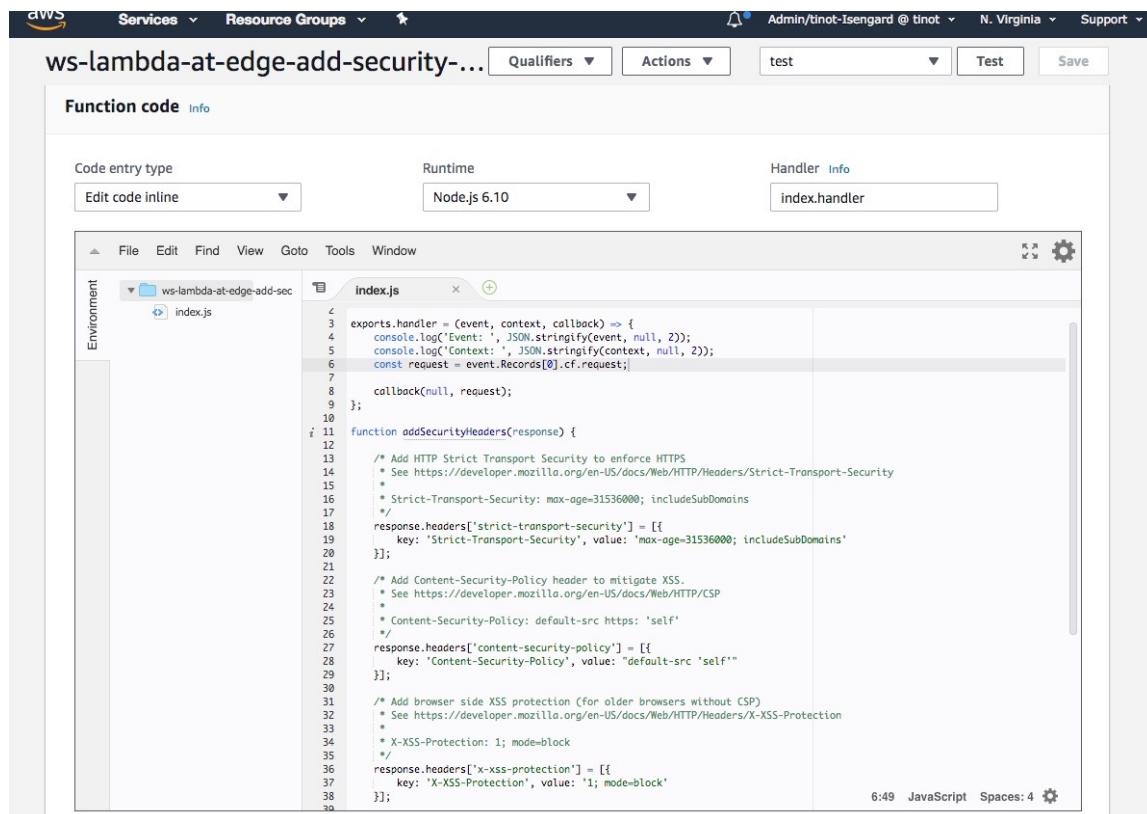
Lambda &
Cloud9 IDE

NodeJS
6.10 & 8.10

CloudWatch
Logs &
Metrics

IAM
integration

Versioning



The screenshot shows the AWS Lambda@Edge function editor interface. The top navigation bar includes 'Services', 'Resource Groups', 'Actions', 'test', 'Test', and 'Save' buttons. The main title is 'ws-lambda-at-edge-add-security-...'. The 'Function code' tab is selected, showing the following code in the 'index.js' file:

```
exports.handler = (event, context, callback) => {
    console.log('Event: ', JSON.stringify(event, null, 2));
    console.log('Context: ', JSON.stringify(context, null, 2));
    const request = event.Records[0].cf.request;
    callback(null, request);
};

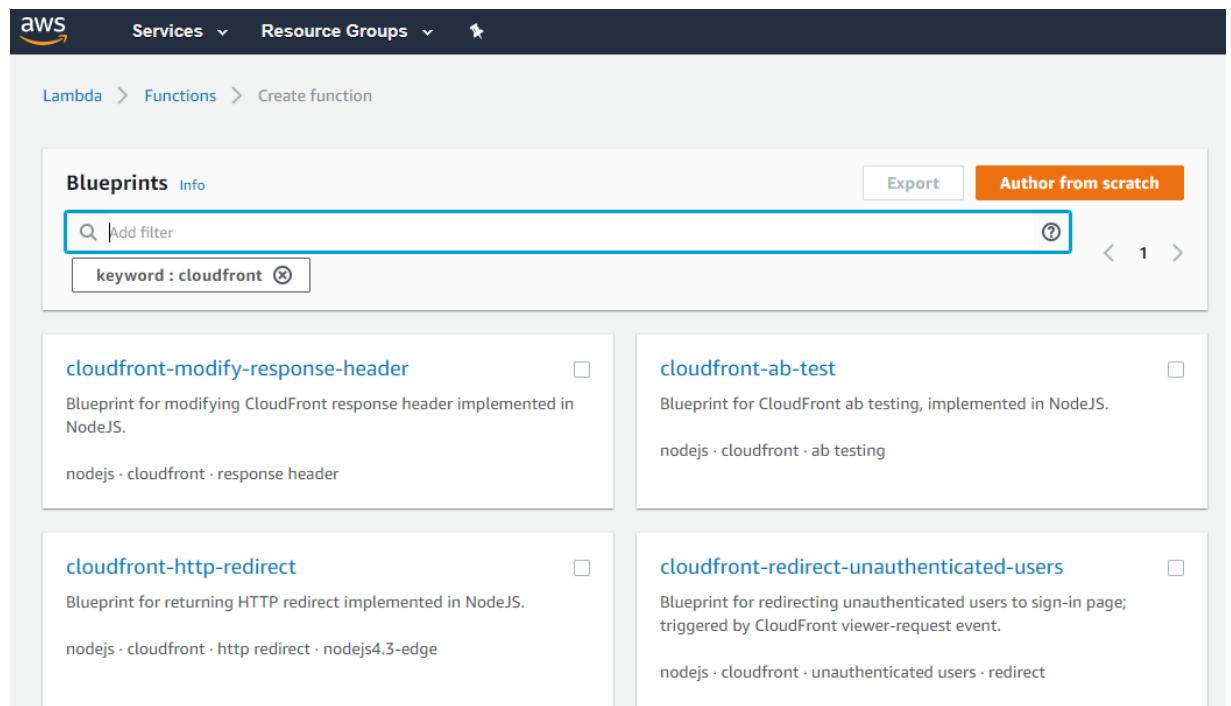
function addSecurityHeaders(response) {
    /* Add HTTP Strict Transport Security to enforce HTTPS
     * See https://developer.mozilla.org/en-US/docs/Web/HTTP/Headers/Strict-Transport-Security
     */
    response.headers['strict-transport-security'] = [
        {key: 'Strict-Transport-Security', value: 'max-age=31536000; includeSubDomains'}
    ];
    /* Add Content-Security-Policy header to mitigate XSS.
     * See https://developer.mozilla.org/en-US/docs/Web/HTTP/CSP
     */
    response.headers['content-security-policy'] = [
        {key: 'Content-Security-Policy', value: "default-src 'self'"}
    ];
    /* Add browser side XSS protection (for older browsers without CSP)
     * See https://developer.mozilla.org/en-US/docs/Web/HTTP/Headers/X-XSS-Protection
     */
    response.headers['x-xss-protection'] = [
        {key: 'X-XSS-Protection', value: '1; mode=block'}
    ];
}
```

The status bar at the bottom right indicates '6:49 JavaScript Spaces: 4'.



Quick start with blueprints

- Content customization based on requests attributes
- user authentication and session validation
- URL rewrites, pretty URLs
- A/B testing and cookie-based sticky sessions



Function structure

```
const LOCATION = 'http://docs.aws.amazon.com/lambda/latest/dg/lambda-edge.html'

exports.handler = (event, context, callback) => {

    const response = {
        status: '302',
        statusDescription: 'Found',
        headers: {
            location: [
                {
                    key: 'Location',
                    value: LOCATION,
                },
            ],
        },
    };
    callback(null, response);
};
```



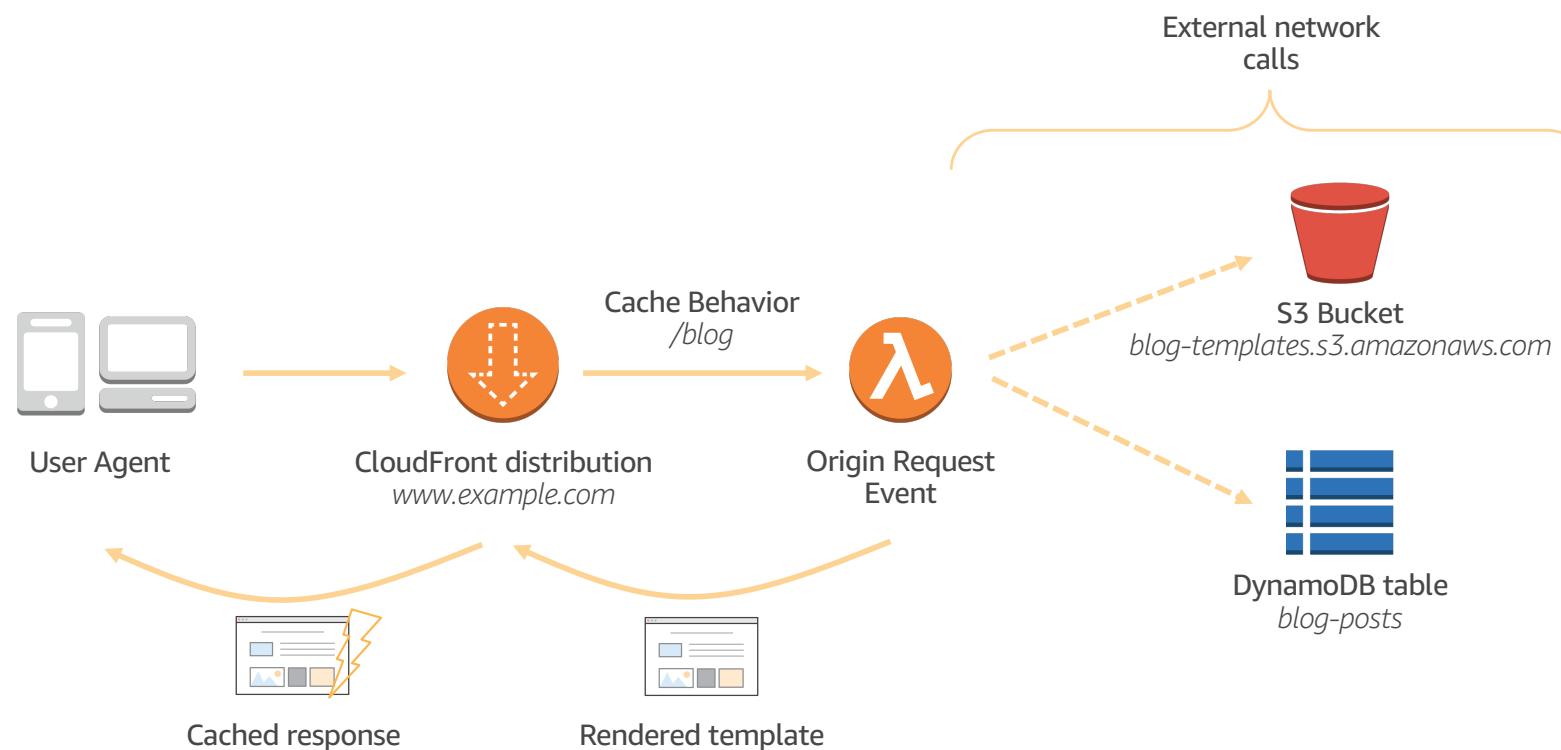
Dynamic content generation (1/3)

```
<h1>{ page.title }</h1>
{{ for section in page.sections }}
<h2>{ section.title }</h2>
<p>{ section.body }</p>
{{ endfor }}
```

```
"page": {
    "title": "Hello",
    "sections": [ {
        "title": "Introduction",
        "body": "The quick..."
    }, { ... } ]
```



Dynamic content generation (2/3)



Dynamic content generation (3/3)

```
const templateBucket = 'blog-templates-123456789012';
const postTable = 'blog-posts';

var AWS = require('aws-sdk');
var Mustache = require('mustache');

var s3 = new AWS.S3({region: 'us-east-1'});
var documentClient = new AWS.DynamoDB.DocumentClient({
  region: 'us-east-1'});

exports.handler = (event, context, callback) => {
  const request = event.Records[0].cf.request;
  const response = {
    status: '200',
    statusDescription: 'OK',
    headers: {
      'cache-control': [
        {key: 'Cache-Control', value: 'max-age=2628000, public'}
      ],
      'content-type': [
        {key: 'Content-Type', value: 'text/html; charset=utf-8'}
      ]},
  };
  callback(null, response);
};
```

```
const ddbParams = {
  TableName: postTable,
  Key: { slug: request['uri'].slice(1) }};

documentClient.get(ddbParams, function(err, resp) {
  if (err) {
    callback(err, null);
    return;
  }

  const template = resp['Item']['template'];
  const data = resp['Item']['data'];
  const s3Params = {
    Bucket: templateBucket,
    Key: template };

  s3.getObject(s3Params, function(err, s3resp) {
    if (err) {
      callback(err, null);
      return;
    }
    const body = s3resp.Body.toString('utf-8');
    response.body = Mustache.render(body, data);
    callback(null, response);
  });
});
```



Security Headers (1/2)

HSTS (tell your viewers to stick to HTTPS)

CORS (allow other domains to access your data relaxing the same-origin policy):

CSP (block content loaded from other sources, reduce XSS risk)

HPKP (pin your public key in the client, prevent MITM attacks with forged certificates):

More:

- X-Frame-Options (prevent clickjacking)
- X-Content-Type-Options (prevent MIME sniffing)
- X-XSS-Protection
- ...

Scan Summary

F

Host:
Scan ID #:
Start Time:
Duration:

Score: 5/100
Tests Passed: 4/11



Scan Summary

A

Host:
Scan ID #:
Start Time:
Duration:

Score: 90/100
Tests Passed: 10/11



Security Headers (2/2)

```
exports.handler = (event, context, callback) => {
  const response = event.Records[0].cf.response;

  // Enforce HTTPS with the HSTS header
  response.headers['strict-transport-security'] = [
    key: 'Strict-Transport-Security', value: 'max-age=31536000; includeSubDomains'
  ];

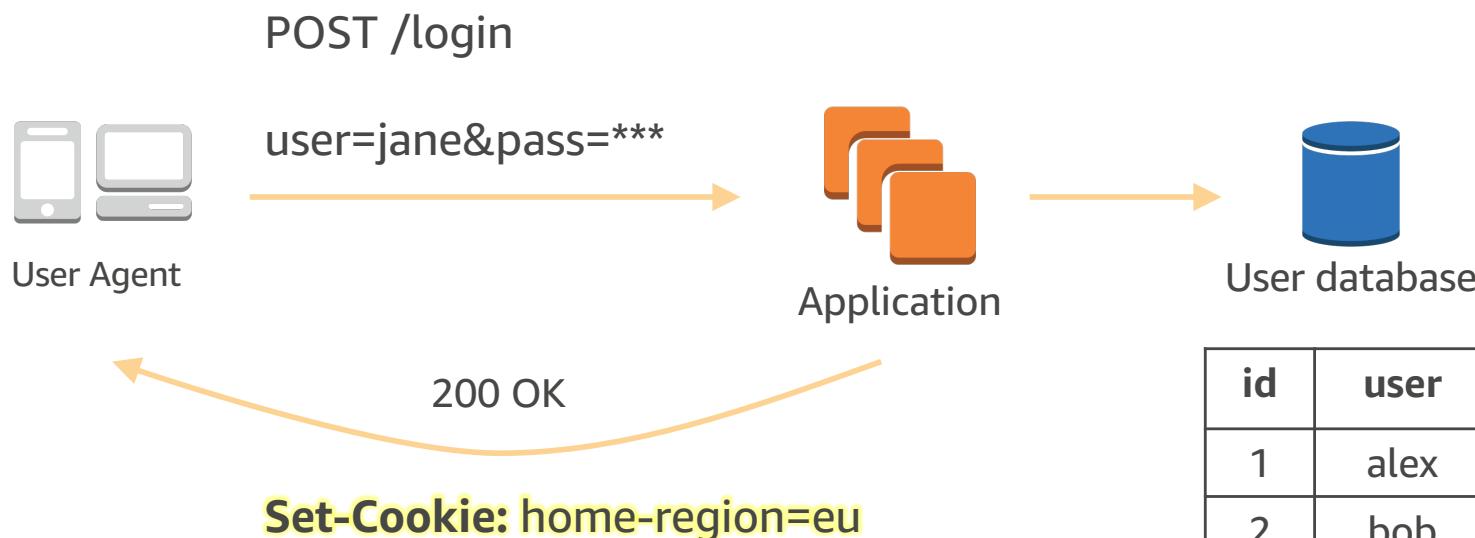
  // Disallow embedded frames
  response.headers['x-frame-options'] = [
    key: 'X-Frame-Options', value: 'DENY'
  ];

  // Remove CORS headers
  delete response.headers['access-control-allow-origin'];
  delete response.headers['access-control-max-age'];
  ...

  callback(null, response);
}
```



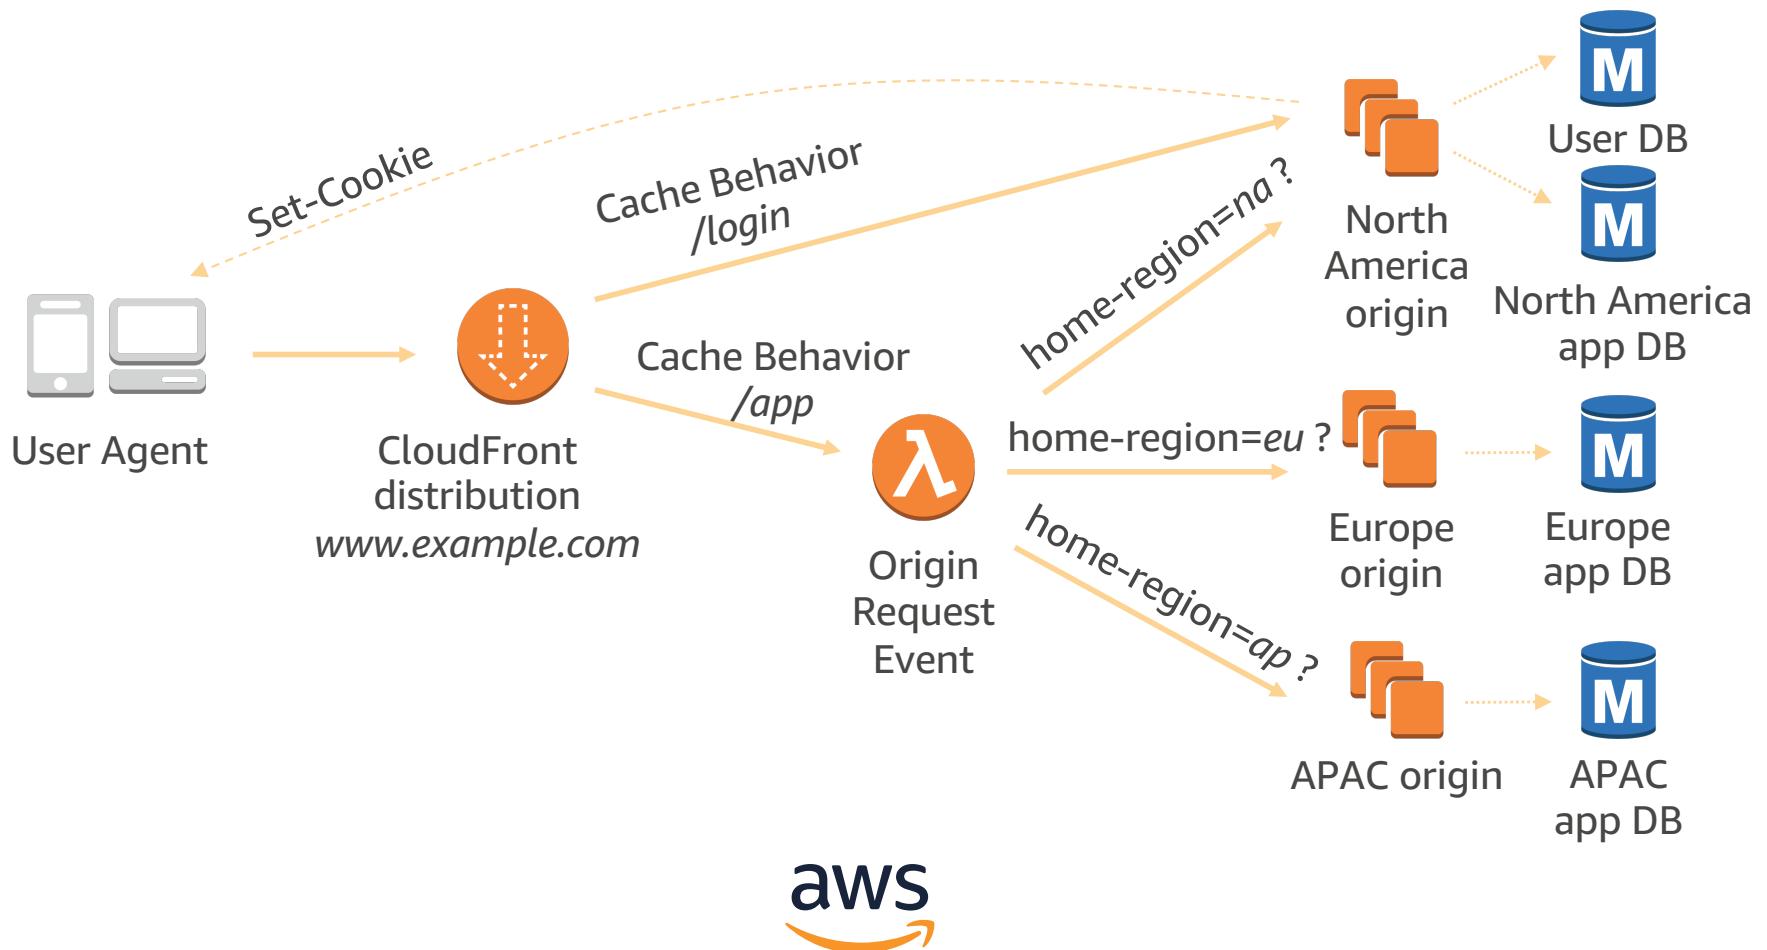
Origin selection (1/2)



id	user	home-region
1	alex	na
2	bob	eu
3	joe	ap
4	jane	eu



Origin selection (1/2)





7 best practices



How much does it cost?

Consider an API with 15M requests/month & a 256MB function executing in 125ms.

Lambda@Edge is charged based on the following two factors:

- Number of requests: $15M * \$0.60/1M = \9
- Memory*Duration resource usage: $15M * 150ms * 256MB * \$0.00005001\$/GBS = 28,1\$$

Viewer Trigger: Total = 37,1\$ per month



#1 Invoke Lambda@Edge only where needed

CloudFront already provide native features:

- **Device identification:** CloudFront-Is-Mobile-Viewer headers
- **Analytics:** Access Logs delivered to S3 & WAF logs
- **Access Control:** signed URLs/Cookies, Geoblocking, WAF

Use the most specific CloudFront behavior:

Cache Behavior Settings

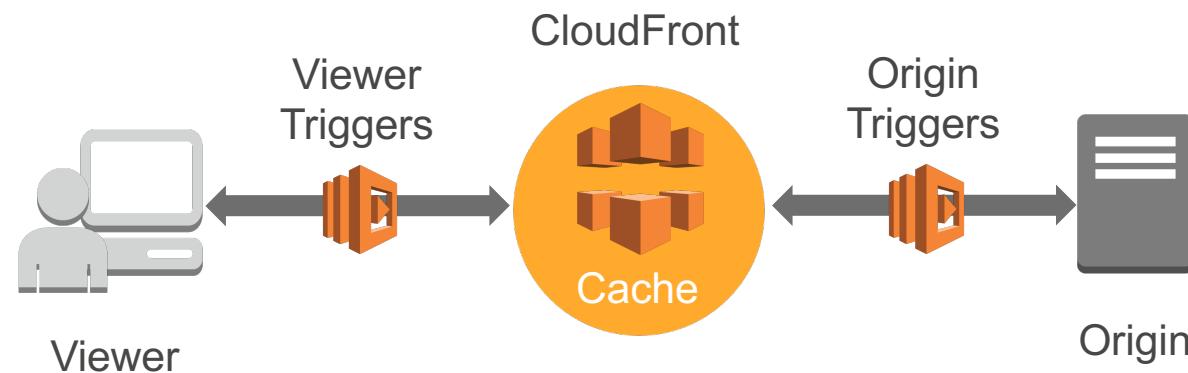
Path Pattern



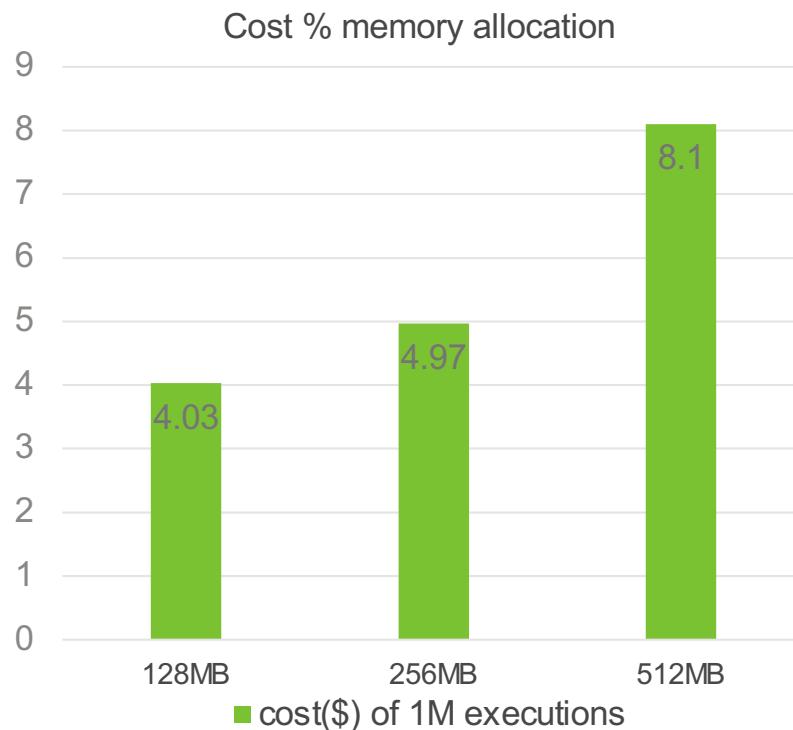
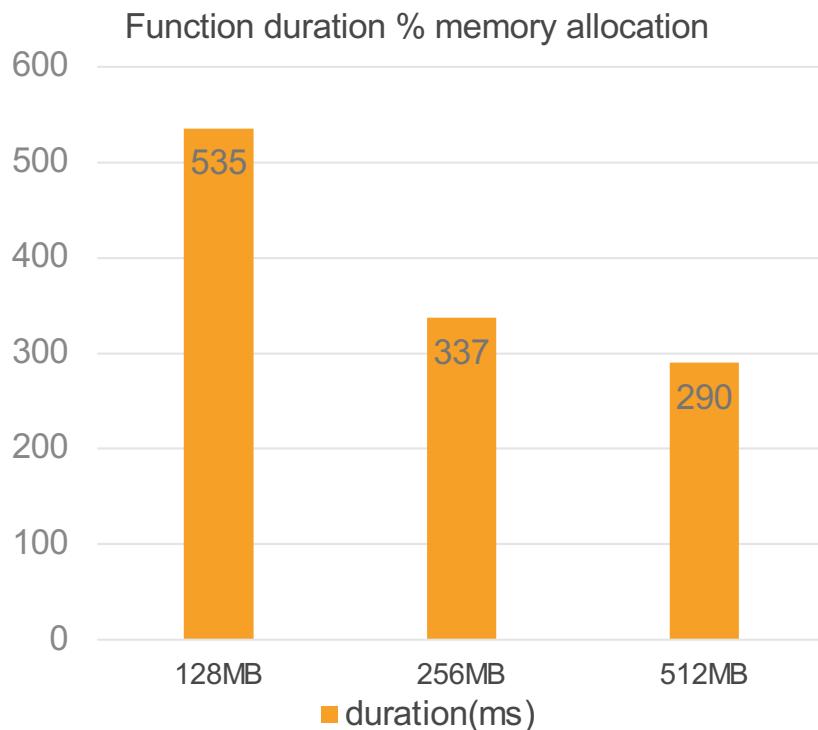
#2 Invoke Lambda@Edge only when needed

Remove it when it's not used any more

For every request or only on cache misses?



#3 Optimize memory configuration



#4 Optimize function code for Node.js

Node.js 8.10 runtime

Reduce deployment package size: external dependencies, lightweight packages, minify, browserfy...

Parallelism using async calls

```
Promise.all([
    httpGet({ hostname: "HTML template", path: ""}),
    ddbGet({ TableName: ddbTableName, Key: { name: "achraf" } })
]).then(responses => { ... }
```



#5 Leverage global variables (1/2)

```
const dns = require("dns");
let bestOrigin;
let expires = 0;

exports.handler = (event, context, callback) => {
    const request = event.Records[0].cf.request;
    getBestOrigin().then((origin) => {
        request.origin.custom.domainName = origin;
        headers.host[0].value = origin;
        callback(null, request);
    });
}
```



#5 Leverage global variables (2/2)

```
function getBestOrigin() {  
    const now = Date.now();  
    if (now < expires) return Promise.resolve(bestOrigin);  
    return new Promise((resolve, reject) => {  
        dns.resolveCname(DNS_HOST, (err, addr) => {  
            bestOrigin = addr[0];  
            expires = now + TTL;  
            resolve(bestOrigin);  
        });  
    });  
}
```



#6 Optimize external network calls

```
const http = require('https');

const keepAliveAgent = new http.Agent({ keepAlive: true, keepAliveMsecs: 2000 });

exports.handler = (event, context, callback) => {

    http.get({ hostname: "api.d.net", path: "/", agent: keepAliveAgent}, (resp) => {

        let data = '';
        resp.on('data', (chunk) => { data += chunk; });
        resp.on('end', () => { resolve(data); });

    });

}

}
```



#7 Know the limits!

Functional:

- Blacklisted/ Read only headers
- Function size – 1MB vs 50MB
- Response size – 40K vs 1MB

Resource allocation

- Memory – 128M vs 3G
- Timeout – 5s vs 30s
- 1K concurrent execution & 10K RPS per region
- Scaling mechanism per region



To learn more about Lambda@Edge

[AWS re:Invent 2017: Taking Serverless to the Edge \(SRV312\)](#)

[AWS re:Invent 2017: Building Serverless Websites with Lambda@Edge \(CTD309\)](#)

[Tutorial: Creating a Simple Lambda@Edge Function](#)

[Testing and Debugging Lambda@Edge Functions](#)

[<https://github.com/aws-samples/aws-lambda-edge-workshops>](#)

[Lambda@Edge Design Best Practices](#)



Some takeaways...

- CloudFront brought storage closer to viewers, Lambda@Edge now brings **compute closer to viewers**.
- Lambda is serverless, Lambda@Edge is **originless**
- Many existing use cases for Lambda@Edge but you have the liberty **to build and innovate**.
- There are things that you can do to **optimize your implementation** in terms of performance and cost.





Thank you!

... time for Q&A

